

## **Microsoft 70-458 Exam Questions & Answers**

Number: 70-458

Passing Score: 700

Time Limit: 180 min

File Version: 38.4



### **Microsoft 70-458 Exam Questions & Answers**

**Exam Name: Transition Your MCTS on SQL Server 2008 to MCSA: SQL Server 2012, Part 2**

## Examsheets

### QUESTION 1

You administer several Microsoft SQL Server 2012 servers. Your company has a number of offices across the world connected by using a wide area network (WAN). Connections between offices vary significantly in both bandwidth and reliability. You need to identify the correct replication method for each scenario. What should you do? (To answer, drag the appropriate replication method or methods to the correct location or locations in the answer area. Each replication method may be used once, more than once, or not at all.)

Replication Method	Scenario
Transactional Replication	Multiple databases on the same low-latency subnet must allow applications to write changes locally, and these changes must be replicated to all related databases.
Peer-to-Peer Replication	An order summary table is repopulated once a week. This table must be replicated to all databases.
Snapshot Replication	Field offices using unreliable connections keep a local copy of the product catalog and process orders locally. These orders must be periodically replicated to all other offices.
Merge Replication	Information in an order-tracking database must be replicated across a low-latency connection as changes occur to multiple reporting databases.

A.

**Correct Answer:**

**Section: (none)**

**Explanation**

**Explanation/Reference:**

Replication Method	Scenario	Replication Method
	Multiple databases on the same low-latency subnet must allow applications to write changes locally, and these changes must be replicated to all related databases	Peer-to-Peer Replication
	An order summary table is repopulated once a week. This table must be replicated to all databases.	Snapshot Replication
	Field offices using unreliable connections keep a local copy of the product catalog and process orders locally. These orders must be periodically replicated to all other offices.	Merge Replication
	Information in an order-tracking database must be replicated across a low-latency connection as changes occur to multiple reporting databases.	Transactional Replication

### QUESTION 2 DRAG DROP

You are a database administrator of a Microsoft SQL Server 2012 environment. The environment contains two servers named SQLServer01 and SQLServer02. The database Contoso exists on SQLServer01. You plan to mirror the Contoso database between SQLServer01 and SQLServer02 by using database mirroring. You need to prepare the Contoso database for database mirroring. Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Back up Contoso on SQLServer01 by using a full backup.

Back up Contoso on SQLServer01 by using a full backup followed by a transaction log backup by using the **NORECOVERY** option.

Back up Contoso on SQLServer01 by using a transaction log backup. Restore the transaction log backup by using the **RECOVERY** option on SQLServer02.

Back up Contoso on SQLServer01 by using a transaction log backup. Restore the transaction log backup by using the **NORECOVERY** option on SQLServer02.

Restore the full database backup of Contoso by using the **NORECOVERY** option on SQLServer02 as Contoso.

Restore the full database backup of Contoso by using the **RECOVERY** option on SQLServer02 as Contoso\_Mirror.



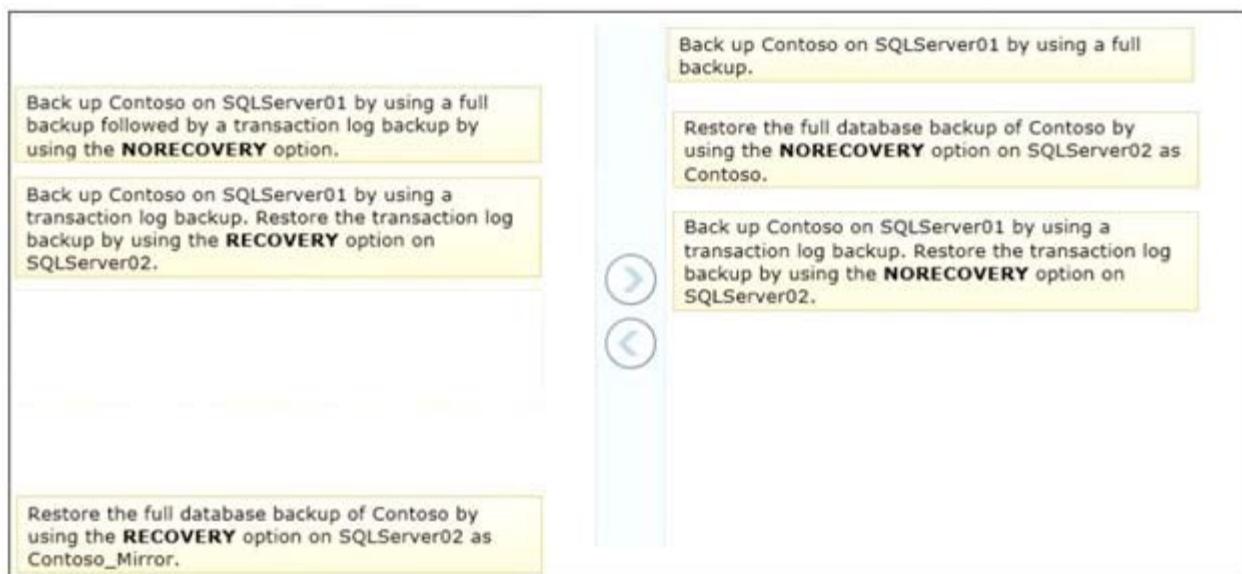
A.

**Correct Answer:**

**Section: (none)**

**Explanation**

**Explanation/Reference:**



### QUESTION 3

You create an availability group that has replicas named HA/Server01 and HA/Server02. Currently, HA/Server01 is the primary replica. You have multiple queries that read data and produce reports from the database. You need to offload the reporting workload to the secondary replica when HA/Server01 is the primary replica. What should you do?

- A. Set the Availability Mode property of HA/Server02 to Asynchronous commit.
- B. Set the Readable Secondary property of HA/Server02 to Read-intent only.
- C. Set the Connections in Primary Role property of HA/Server01 to Allow read/write connections.
- D. Set the Availability Mode property of HA/Server01 to Asynchronous commit.

**Correct Answer:** C

**Section:** (none)

**Explanation**

**Explanation/Reference:**

### QUESTION 4

You administer several Microsoft SQL Server 2012 database servers. Merge replication has been configured for an application that is distributed across offices throughout a wide area network (WAN). Many of the tables involved in replication use the XML and varchar (max) data types. Occasionally, merge replication fails due to timeout errors. You need to reduce the occurrence of these timeout errors. What should you do?

- A. Set the Merge agent on the problem subscribers to use the slow link agent profile.
- B. Create a snapshot publication, and reconfigure the problem subscribers to use the snapshot publication.
- C. Change the Merge agent on the problem subscribers to run continuously.
- D. Set the Remote Connection Timeout on the Publisher to 0.

**Correct Answer:** A

**Section:** (none)

**Explanation**

**Explanation/Reference:**

**QUESTION 5**

You administer two Microsoft SQL Server 2012 servers. Each server resides in a different, untrusted domain. You plan to configure database mirroring. You need to be able to create database mirroring endpoints on both servers. What should you do?

- A. Configure the SQL Server service account to use Network Service.
- B. Use a server certificate.
- C. Use a database certificate.
- D. Configure the SQL Server service account to use Local System.

**Correct Answer:** B

**Section:** (none)

**Explanation**

**Explanation/Reference:**

**QUESTION 6**

You are migrating a database named Orders to a new server that runs Microsoft SQL Server 2012. You attempt to add the [Corpnet\User1] login to the database. However, you receive the following error message: "User already exists in current database."

You need to configure the [Corpnet\User1] login to be able to access the Orders database and retain the original permissions. You need to achieve this goal by using the minimum required permissions. Which Transact-SQL statement should you use?

- A. DROP USER [User1];  
CREATE USER [Corpnet\User1] FOR LOGIN [Corpnet\User1]; ALTER ROLE [db\_owner] ADD MEMBER [Corpnet\User1];
- B. ALTER SERVER RCLSID sysadmin] ADD MEMBER [Corpnet\User1];
- C. ALTER USER [Corpnet\User1] WITH LOGIN [Corpnet\User1];
- D. ALTER ROLE [db owner] ADD MEMBER [Corpnet\User1];

**Correct Answer:** C

**Section:** (none)

**Explanation**

**Explanation/Reference:**

**QUESTION 7**

You administer a Microsoft SQL Server 2012 database. You configure Transparent Data Encryption (TDE) on the Orders database by using the following statements:

```
CREATE MASTER KEY ENCRYPTION BY PASSWORD = 'MyPassword1!';
CREATE CERTIFICATE TDE_Certificate WITH SUBJECT = 'TDE Certificate';

BACKUP CERTIFICATE TDE_Certificate TO FILE = 'd:\TDE_Certificate.cer'
WITH PRIVATE KEY (FILE = 'd:\TDE_Certificate.key', ENCRYPTION BY PASSWOR

CREATE DATABASE ENCRYPTION KEY
WITH ALGORITHM = AES_256
ENCRYPTION BY SERVER CERTIFICATE TDE_Certificate;

ALTER DATABASE Orders SET ENCRYPTION ON;
```

You attempt to restore the Orders database and the restore fails. You copy the encryption file to the original location. A hardware failure occurs and so a new server must be installed and configured. After installing SQL Server to the new server, you restore the Orders database and copy the encryption files to their original location. However, you are unable to access the database. You need to be able to restore the database. Which Transact-SQL statement should you use before attempting the restore?

- A. CREATE ASSEMBLY TDE\_Assembly  
    FROM 'd:\TDE\_Certificate.cer'  
    WITH PERMISSION\_SET = SAFE;  
    GO  
    CREATE CERTIFICATE TDE\_Certificate FROM ASSEMBLY TDE\_Assembly;
  - B. CREATE CERTIFICATE TDE\_Certificate FROM EXECUTABLE FILE = 'd:\TDE\_
  - C. CREATE CERTIFICATE TDE\_Certificate FROM FILE = 'd:\TDE\_Certificate'  
    WITH PRIVATE KEY (FILE = 'd:\TDE\_Certificate.key', DECRYPTION BY P
  - D. DECLARE @startdate date  
    SET @startdate = GETDATE()  
    CREATE CERTIFICATE TDE\_Certificate FROM FILE = 'd:\TDE\_Certificate'  
    WITH START\_DATE = @startdate;
- A. Option A  
B. Option B  
C. Option C  
D. Option D

**Correct Answer: C**

**Section: (none)**

**Explanation**

**Explanation/Reference:**

**QUESTION 8**

You administer a Microsoft SQL Server 2012 database that has Trustworthy set to On. You create a stored procedure that returns database-level information from Dynamic Management Views. You grant User1 access to execute the stored procedure. You need to ensure that the stored procedure returns the required information when User1 executes the stored procedure. You need to achieve this goal by granting the minimum permissions required. What should you do? (Each correct answer presents a complete solution. Choose all that apply.)

- A. Create a SQL Server login that has VIEW SERVER STATE permissions. Create an application role and a secured password for the role.
- B. Modify the stored procedure to include the EXECUTE AS OWNER statement. Grant VIEW SERVER STATE permissions to the owner of the stored procedure.
- C. Create a SQL Server login that has VIEW SERVER STATE permissions. Modify the stored procedure to include the EXECUTE AS {newlogin} statement.
- D. Grant the db\_owner role on the database to User1.
- E. Grant the sysadmin role on the database to User1.

**Correct Answer:** DE

**Section:** (none)

**Explanation**

**Explanation/Reference:**

**QUESTION 9**

You are the lead database administrator (DBA) of a Microsoft SQL Server 2012 environment. All DBAs are members of the DOMAIN\JrDBAs Active Directory group. You grant DOMAIN\JrDBAs access to the SQL Server. You need to create a server role named SpecialDBARole that can perform the following functions:

- View all databases.
- View the server state.
- Assign GRANT, DENY, and REVOKE permissions on logins. You need to add DOMAIN\JrDBAs to the server role. You also need to provide the least level of privileges necessary. Which SQL statement or statements should you use? Choose all that apply.

- A. CREATE SERVER ROLE [SpecialDBARole] AUTHORIZATION setupadmin;
- B. ALTER SERVER ROLE [SpecialDBARole] ADD MEMBER [DOMAIN\JrDBAs];
- C. CREATE SERVER ROLE [SpecialDBARole] AUTHORIZATION securityadmin;
- D. GRANT VIEW DEFINITION TO [SpecialDBARole];
- E. CREATE SERVER ROLE [SpecialDBARole] AUTHORIZATION serveradmin;
- F. GRANT VIEW SERVER STATE, VIEW ANY DATABASE TO [SpecialDBARole];

**Correct Answer:** CEF

**Section:** (none)

**Explanation**

**Explanation/Reference:**

**QUESTION 10**

**DRAG DROP**

You administer a Microsoft SQL Server 2012 instance that contains a database of confidential data. You need to encrypt the database files at the page level. You also need to encrypt the transaction log files. Which four actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to

the answer area and arrange them in the correct order.)



A.

**Correct Answer:**

**Section: (none)**

**Explanation**

**Explanation/Reference:**



#### **QUESTION 11**

You administer a Microsoft SQL Server 2012 database. The database contains a table that has the following definition:

```
CREATE TABLE [Sales].[Customer] (
    [CustomerID] int NOT NULL,
    [CustomerName] nvarchar(50) NOT NULL,
    [TerritoryID] int NULL,
    [LastContactDate] datetimeoffset NULL,
    [CustomerType] nchar(1) NOT NULL,
    [Notes] varchar(250) NULL
)
```

You want to export data from the table to a flat file by using the SQL Server Import and Export Wizard. You need to ensure that the following requirements are met:

\*The first row of the file contains the first row of data.

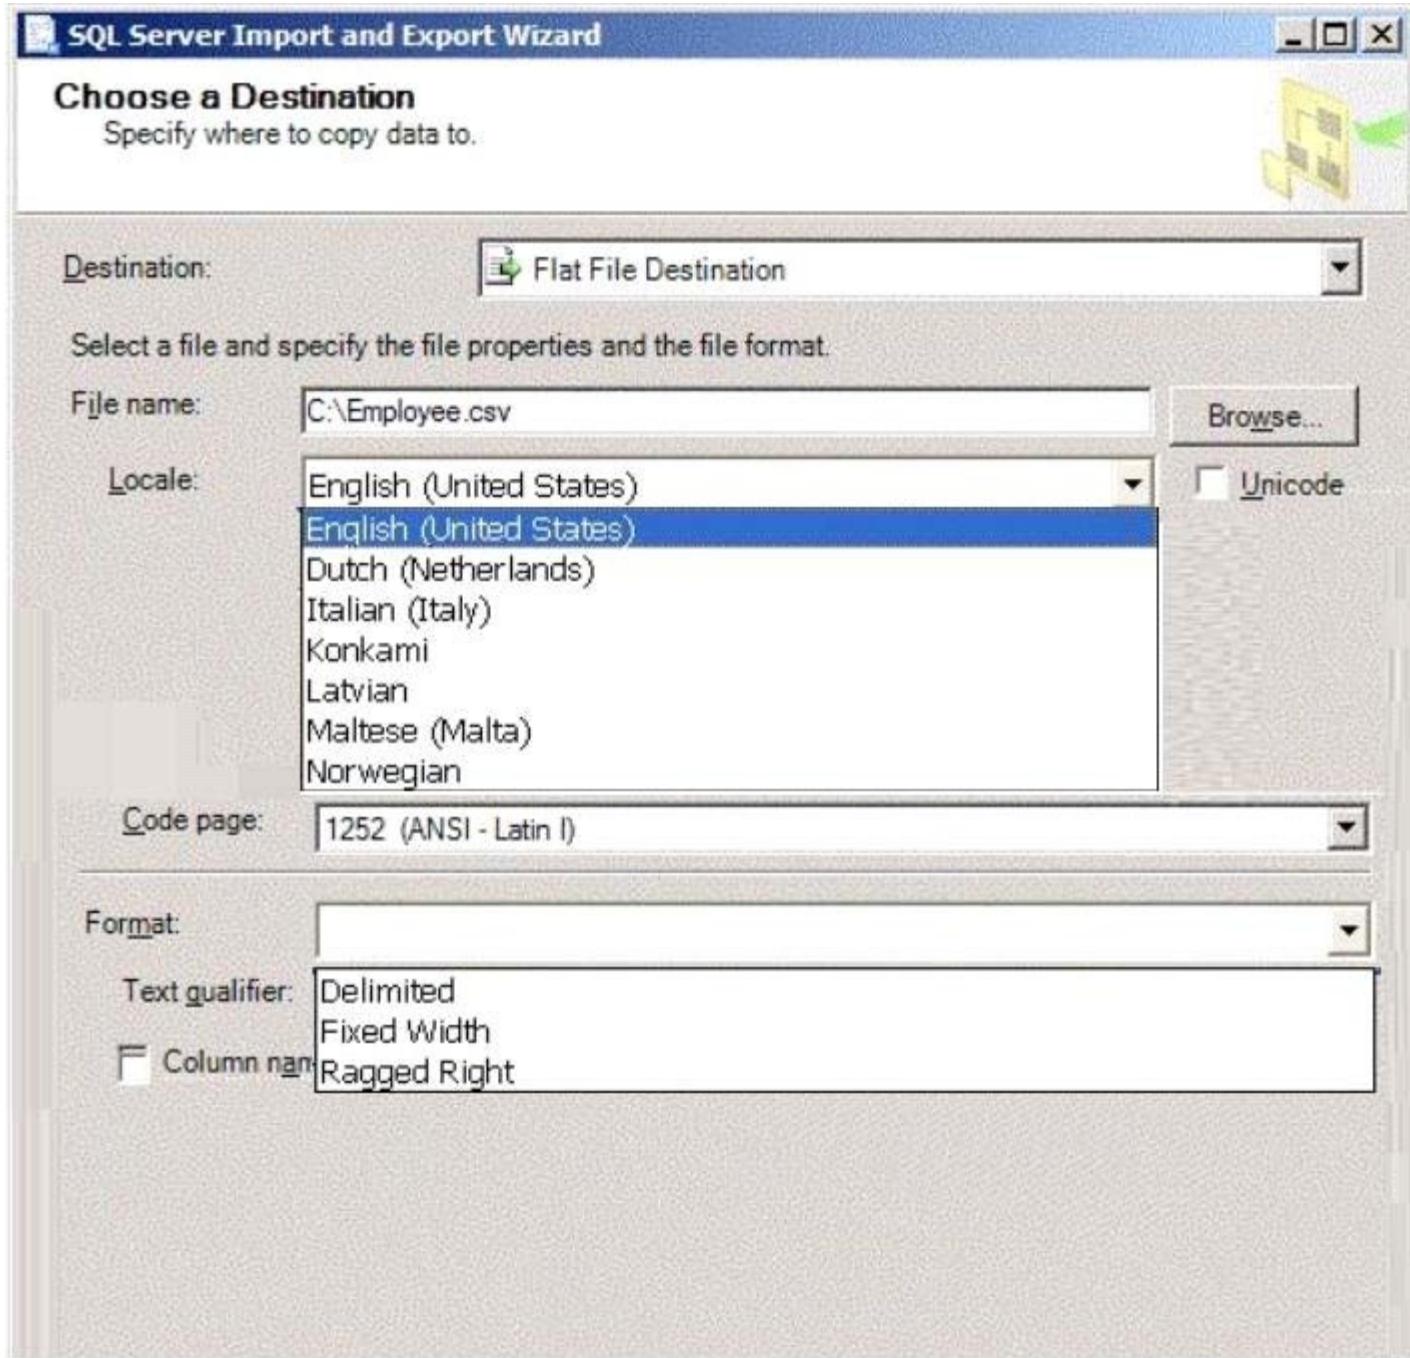
\*Each record is of the same length.

\*The date follows the U.S. date format.

\*The file supports international characters.

What should you do? (To answer, configure the appropriate option or options in the dialog box in the answer area.)





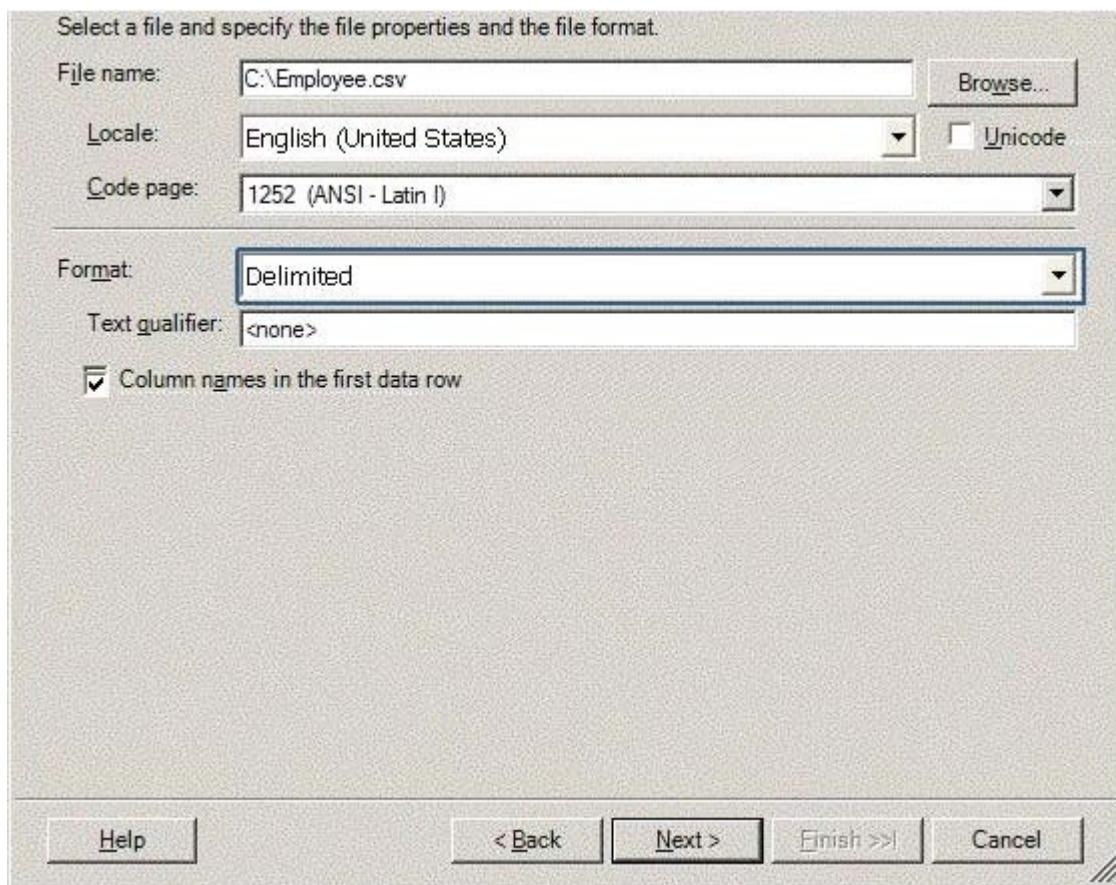
A.

**Correct Answer:**

**Section:** (none)

**Explanation**

**Explanation/Reference:**



### QUESTION 12

DRAG DROP

You administer a Microsoft SQL Server 2012 database. You use an OrderDetail table that has the following definition:

```
CREATE TABLE [dbo].[OrderDetail]
([SalesOrderID] [int] NOT NULL,
 [SalesOrderDetailID] [int] IDENTITY(1,1) NOT NULL,
 [CarrierTrackingNumber] [nvarchar](25) NULL,
 [OrderQty] [smallint] NOT NULL,
 [ProductID] [int] NOT NULL,
 [SpecialOfferID] [int] NULL,
 [UnitPrice] [money] NOT NULL);
```

You need to create a non-clustered index on the SalesOrderID column in the OrderDetail table to include only rows that contain a value in the CarrierTrackingNumber column. Which four Transact-SQL statements should you use? (To answer, move the appropriate statements from the list of statements to the answer area and arrange them in the correct order.)

WHERE	
FILTER ON	
CarrierTrackingNumber IS NOT NULL;	
ON dbo.OrderDetail (SalesOrderID)	
ON dbo.OrderDetail (SalesOrderID) AS FILTERED_INDEX	
CREATE NONCLUSTERED INDEX FIndx_CarrierTrackingNumber	
CREATE NONCLUSTERED FILTERED INDEX FIndx_CarrierTrackingNumber	

A.

**Correct Answer:**

**Section: (none)**

**Explanation**

**Explanation/Reference:**

FILTER ON	CREATE NONCLUSTERED INDEX FIndx_CarrierTrackingNumber
ON dbo.OrderDetail (SalesOrderID)	ON dbo.OrderDetail (SalesOrderID)
WHERE	WHERE
CarrierTrackingNumber IS NOT NULL;	CarrierTrackingNumber IS NOT NULL;
ON dbo.OrderDetail (SalesOrderID) AS FILTERED_INDEX	
CREATE NONCLUSTERED FILTERED INDEX FIndx_CarrierTrackingNumber	

### QUESTION 13

You administer a SQL 2012 server that contains a database named SalesDb. SalesDb contains a schema named Customers that has a table named Regions. A user named UserA is a member of a role named Sales. UserA is granted the Select permission on the Regions table. The Sales role is granted the Select permission on the Customers schema. You need to remove the Select permission for UserA on the Regions table. You also need to ensure that UserA can still access all the tables in the Customers schema, including the Regions table, through the Sales role permissions. Which Transact-SQL statement should you use?

- A. DENY SELECT ON Object::Regions FROM UserA
- B. DENY SELECT ON Schema::Customers FROM UserA
- C. EXEC sp\_addrolemember 'Sales', 'UserA'
- D. REVOKE SELECT ON Object:: Regions FROM UserA
- E. REVOKE SELECT ON Object::Regions FROM Sales
- F. EXEC sp\_droprolemember 'Sales', 'UserA'

- G. REVOKE SELECT ON Schema::Customers FROM UserA
- H. DENY SELECT ON Object::Regions FROM Sales
- I. DENY SELECT ON Schema:: Customers FROM Sales
- J. REVOKE SELECT ON Schema:: Customers FROM Sales

**Correct Answer:** D

**Section:** (none)

**Explanation**

**Explanation/Reference:**

Get Complete Collection of 70-458 Exam's Question and Answers.  
<http://www.Test4Prep.com>

# Test4Prep

## QUESTION 14

You administer a Microsoft SQL Server 2012 instance that contains a financial database hosted on a storage area network (SAN). The financial database has the following characteristics:

- A data file of 2 terabytes is located on a dedicated LUN (drive D).
- A transaction log of 10 GB is located on a dedicated LUN (drive E).
- Drive D has 1 terabyte of free disk space.
- Drive E has 5 GB of free disk space.

The database is continually modified by users during business hours from Monday through Friday between 09:00 hours and 17:00 hours. Five percent of the existing data is modified each day. The Finance department loads large CSV files into a number of tables each business day at 11:15 hours and 15:15 hours by using the BCP or BULK INSERT commands. Each data load adds 3 GB of data to the database. These data load operations must occur in the minimum amount of time. A full database backup is performed every Sunday at 10:00 hours. Backup operations will be performed every two hours (11:00, 13:00, 15:00, and 17:00) during business hours. You need to ensure that your backup will continue if any invalid checksum is encountered. Which backup option should you use?

- A. STANDBY
- B. Differential
- C. FULL
- D. CHECKSUM
- E. BULK\_LOGGED
- F. CONTINUE\_AFTER\_ERROR
- G. SIMPLE
- H. DBO\_ONLY
- I. COPY\_ONLY
- J. SKIP
- K. RESTART
- L. Transaction log
- M. NO\_CHECKSUM
- N. NORECOVERY

**Correct Answer:** F

**Section: (none)****Explanation****Explanation/Reference:****QUESTION 15**

You administer a Microsoft SQL Server 2012 instance that contains a financial database hosted on a storage area network (SAN). The financial database has the following characteristics:

- A data file of 2 terabytes is located on a dedicated LUN (drive D).
- A transaction log of 10 GB is located on a dedicated LUN (drive E).
- Drive D has 1 terabyte of free disk space.
- Drive E has 5 GB of free disk space.

The database is continually modified by users during business hours from Monday through Friday between 09:00 hours and 17:00 hours. Five percent of the existing data is modified each day. The Finance department loads large CSV files into a number of tables each business day at 11:15 hours and 15:15 hours by using the BCP or BULK INSERT commands. Each data load adds 3 GB of data to the database. These data load operations must occur in the minimum amount of time. A full database backup is performed every Sunday at 10:00 hours. Backup operations will be performed every two hours (11:00, 13:00, 15:00, and 17:00) during business hours. On Wednesday at 10:00 hours, the development team requests you to refresh the database on a development server by using the most recent version. You need to perform a full database backup that will be restored on the development server. Which backup option should you use?

- A. NORECOVERY
- B. FULL
- C. NO\_CHECKSUM
- D. CHECKSUM
- E. Differential
- F. BULK\_LOGGED
- G. STANDBY
- H. RESTART
- I. SKIP
- J. Transaction log
- K. DBO ONLY
- L. COPY\_ONLY
- M. SIMPLE
- N. CONTINUE\_AFTER\_ERROR

**Correct Answer: K****Section: (none)****Explanation****Explanation/Reference:****QUESTION 16**

You administer a Microsoft SQL Server 2012 instance that contains a financial database hosted on a storage area network (SAN). The financial database has the following characteristics:

- A data file of 2 terabytes is located on a dedicated LUN (drive D).
- A transaction log of 10 GB is located on a dedicated LUN (drive E).
- Drive D has 1 terabyte of free disk space.
- Drive E has 5 GB of free disk space.

The database is continually modified by users during business hours from Monday through Friday between 09:00 hours and 17:00 hours. Five percent of the existing data is modified each day. The Finance department loads large CSV files into a number of tables each business day at 11:15 hours and 15:15 hours by using the

BCP or BULK INSERT commands. Each data load adds 3 GB of data to the database. These data load operations must occur in the minimum amount of time. A full database backup is performed every Sunday at 10:00 hours. Backup operations will be performed every two hours (11:00, 13:00, 15:00, and 17:00) during business hours. You implement log shipping of the financial database to another SQL Server 2012 instance. You decide to failover to this secondary database. You need to ensure that all transactions will be replicated to the secondary database. Which backup option should you use?

- A. Differential
- B. Transaction
- C. FULL
- D. SIMPLE
  - E. BULK\_LOGGED
- E. SKIP
- F. RESTART
- G. STANDBY
- H. CHECKSUM
- I. DBO\_ONLY
- J. COPY ONLY
- K. NORECOVERY
- L. NO\_CHECKSUM
- M. CONTINUE\_AFTER\_ERROR

**Correct Answer:** L

**Section:** (none)

**Explanation**

**Explanation/Reference:**

### **QUESTION 17**

You are creating a SQL Server Integration Services (SSIS) package to retrieve product data from two different sources. One source is hosted in a SQL Azure database. Each source contains products for different distributors. Products for each distributor source must be combined for insertion into a single product table destination. You need to select the appropriate data flow transformation to meet this requirement. Which transformation types should you use? (Each correct answer presents a complete solution. Choose all that apply.)

- A. Multicast
- B. Merge Join
- C. Term Extraction
- D. union All
- E. Merge

**Correct Answer:** AE

**Section:** (none)

**Explanation**

**Explanation/Reference:**

### **QUESTION 18**

You are designing a SQL Server Integration Services (SSIS) package that uses the Fuzzy Lookup transformation. The reference data to be used in the transformation does not change. You need to reuse the Fuzzy Lookup match index to increase performance and reduce maintenance. What should you do?

- A. Select the GenerateAndPersistNewIndex option in the Fuzzy Lookup Transformation Editor.
- B. Select the GenerateNewIndex option in the Fuzzy Lookup Transformation Editor.
- C. Select the DropExistingMatchIndex option in the Fuzzy Lookup Transformation Editor.
- D. Execute the sp\_FuzzyLookupTableMaintenanceUninstall stored procedure.
  - E Execute the sp\_FuzzyLookupTableMaintenanceInvoke stored procedure.

**Correct Answer:** A

**Section:** (none)

**Explanation**

**Explanation/Reference:**

#### **QUESTION 19**

You are creating a SQL Server Master Data Services (MDS) model for a company. The source data for the company is stored in a single table that contains the manager-to-subordinate relationships. You need to create a hierarchy representing the organizational structure of the company. Which hierarchy type should you use?

- A. Natural
- B. Explicit
- C. Parent
- D. Recursive

**Correct Answer:** D

**Section:** (none)

**Explanation**

**Explanation/Reference:**

#### **QUESTION 20**

You are the data steward for a Business Intelligence project. You must identify duplicate rows stored in a SQL Server table and output discoveries to a CSV file. A Data Quality Services (DQS) knowledge base has been created to support this project. You need to produce the CSV file with the least amount of development effort. What should you do?

- A. Create an Integration Services package and use a Data Profiling transform.
- B. Create a custom .NET application based on the Knowledgebase class.
  - C Create a data quality project.
- C. Create a CLR stored procedure based on the Knowledgebase class.
- D. Create a Master Data Services (MDS) business rule.

**Correct Answer:** C

**Section:** (none)

**Explanation**

**Explanation/Reference:**

#### **QUESTION 21**

You are implementing the indexing strategy for a fact table in a data warehouse. The fact table is named Quotes. The table has no indexes and consists of seven columns:

- [ID]
- [QuoteDate]

- [Open]
- [Close]
- [High]
- [Low]
- [Volume]

Each of the following queries must be able to use a columnstore index:

- SELECT AVG ([Close]) AS [AverageClose] FROM Quotes WHERE [QuoteDate] BETWEEN '20100101' AND '20101231'.
- SELECT AVG([High] - [Low]) AS [AverageRange] FROM Quotes WHERE [QuoteDate] BETWEEN '20100101' AND '20101231'.
- SELECT SUM([Volume]) AS [SumVolume] FROM Quotes WHERE [QuoteDate] BETWEEN '20100101' AND '20101231'.

You need to ensure that the indexing strategy meets the requirements. The strategy must also minimize the number and size of the indexes.

What should you do?

- Create one columnstore index that contains [ID], [Close], [High], [Low], [Volume], and [QuoteDate].
- Create three columnstore indexes:
  - One containing [QuoteDate] and [Close]
  - One containing [QuoteDate], [High], and [Low]
  - One containing [QuoteDate] and [Volume]
- Create one columnstore index that contains [QuoteDate], [Close], [High], [Low], and [Volume].
- Create two columnstore indexes:
  - One containing [ID], [QuoteDate], [Volume], and [Close]
  - One containing [ID], [QuoteDate], [High], and [Low]

**Correct Answer: C**

**Section: (none)**

**Explanation**

**Explanation/Reference:**

### QUESTION 22

You are designing a data warehouse with two fact tables. The first table contains sales per month and the second table contains orders per day. Referential integrity must be enforced declaratively. You need to design a solution that can join a single time dimension to both fact tables. What should you do?

- Join the two fact tables.
- Merge the fact tables.
- Create a time dimension that can join to both fact tables at their respective granularity.
- Create a surrogate key for the time dimension.

**Correct Answer: C**

**Section: (none)**

**Explanation**

**Explanation/Reference:**

### QUESTION 23

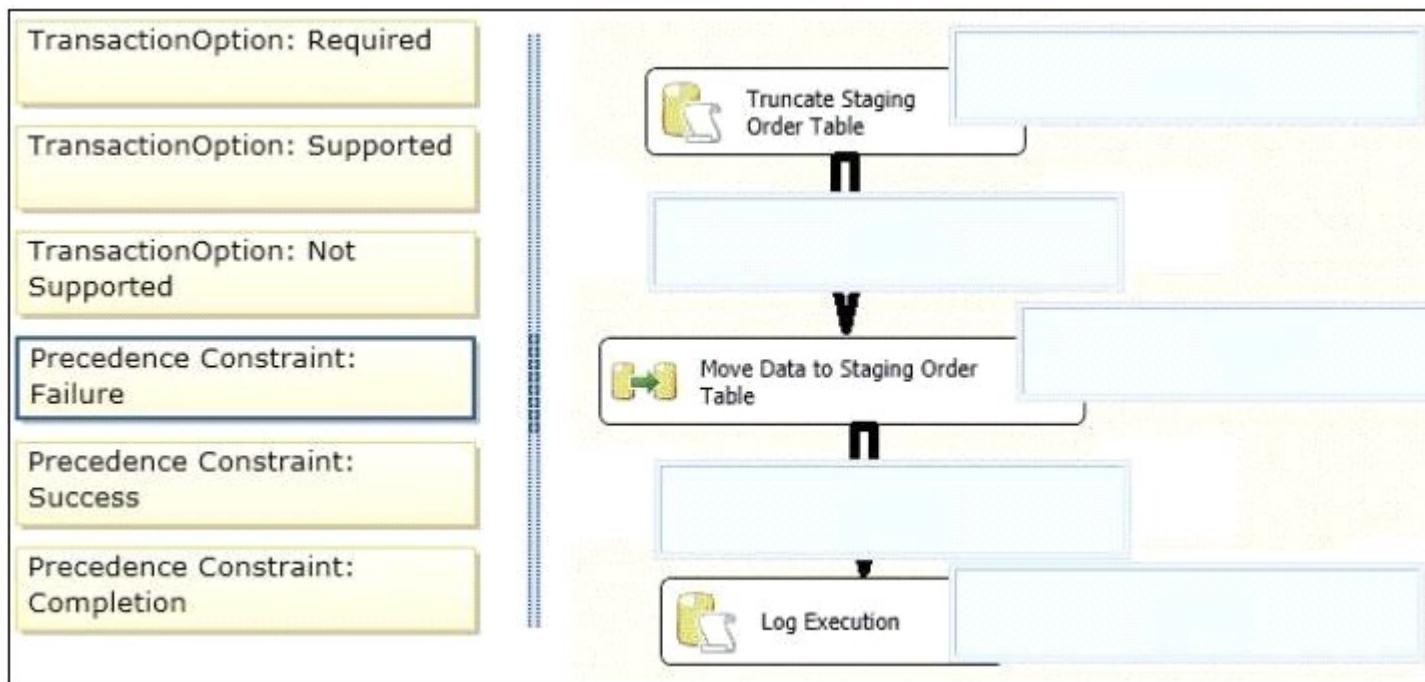
**DRAG DROP**

You are designing a SQL Server Integration Services (SSIS) package. The package moves order-related data to a staging table named Order. Every night the staging data is truncated and then all the recent orders from the online store database are inserted into the staging table. Your package must meet the following

requirements:

- If the truncate operation fails, the package execution must stop and report an error.
- If the Data Flow task that moves the data to the staging table fails, the entire refresh operation must be rolled back.
- For auditing purposes, a log entry must be entered in a SQL log table after each execution of the Data Flow task.

The TransactionOption property for the package is set to Required. You need to design the package to meet the requirements. How should you design the control flow for the package? (To answer, drag the appropriate setting from the list of settings to the correct location or locations in the answer area.)



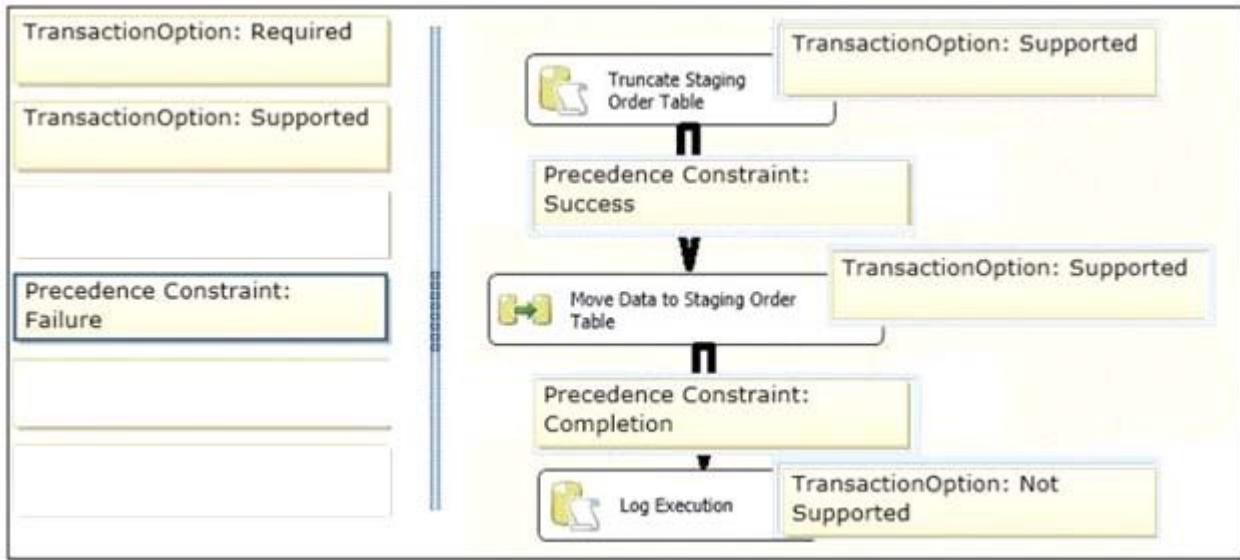
A.

**Correct Answer:**

**Section: (none)**

**Explanation**

**Explanation/Reference:**



## QUESTION 24

### DRAG DROP

You develop a SQL Server Integration Services (SSIS) project by using the Project Deployment model. The project contains many packages. It is deployed on a server named Development1. The project will be deployed to several servers that run SQL Server 2012. The project accepts one required parameter. The data type of the parameter is a string. A SQL Agent job is created that will call the master.dtsx package in the project. A job step is created for the SSIS package. The job must pass the value of an SSIS Environment Variable to the project parameter. The value of the Environment Variable must be configured differently on each server that runs SQL Server. The value of the Environment Variable must provide the server name to the project parameter. You need to configure SSIS on the Development1 server to pass the Environment Variable to the package. Which four actions should you perform in sequence by using SQL Server Management Studio? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

In the SSIS catalog, create an Environment named **DevelopmentEnv**.

In the **DevelopmentEnv** Environment, create an Environment Variable. Name the Environment Variable **PackageEnvironment**. Set the Environment Variable value to **Development1**.

In the SQL Agent job, choose the **DevelopmentEnv** Environment in the Configuration tab of the SQL Server Integration Services job step.

In the project configuration, create an Environment reference to **DevelopmentEnv**. Assign the value of the **PackageEnvironment** Environment Variable to the project parameter.

In the project configuration, edit the value of the project parameter to **DevelopmentEnv**.

In the SSIS catalog, create a folder named **DevelopmentEnv**.



A.

**Correct Answer:**

**Section: (none)**

**Explanation**

**Explanation/Reference:**

In the SSIS catalog, create an Environment named **DevelopmentEnv**.

In the SQL Agent job, choose the **DevelopmentEnv** Environment in the Configuration tab of the SQL Server Integration Services job step.

In the SSIS catalog, create a folder named **DevelopmentEnv**.

In the project configuration, create an Environment reference to **DevelopmentEnv**. Assign the value of the **PackageEnvironment** Environment Variable to the project parameter.

In the project configuration, edit the value of the project parameter to **DevelopmentEnv**.

In the **DevelopmentEnv** Environment, create an Environment Variable. Name the Environment Variable **PackageEnvironment**. Set the Environment Variable value to **Development1**.

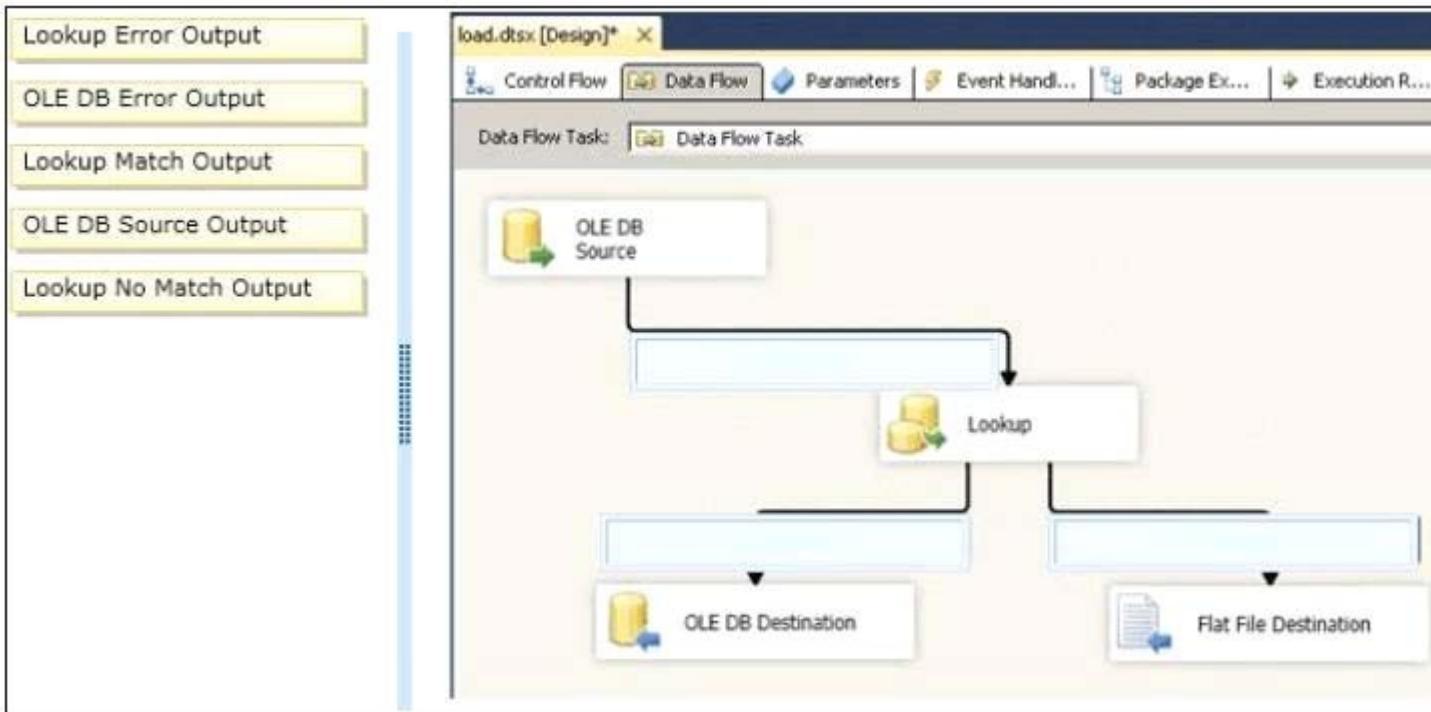
## QUESTION 25

DRAG DROP

You are developing a SQL Server Integration Services (SSIS) package to insert new data into a data mart. The package uses a Lookup transformation to find matches between the source and destination. The data flow has the following requirements:

- New rows must be inserted.
- Lookup failures must be written to a flat file.

In the Lookup transformation, the setting for rows with no matching entries is set to Redirect rows to no match output. You need to configure the package to direct data into the correct destinations. How should you design the data flow outputs? (To answer, drag the appropriate transformation from the list of answer options to the correct location in the answer area.)



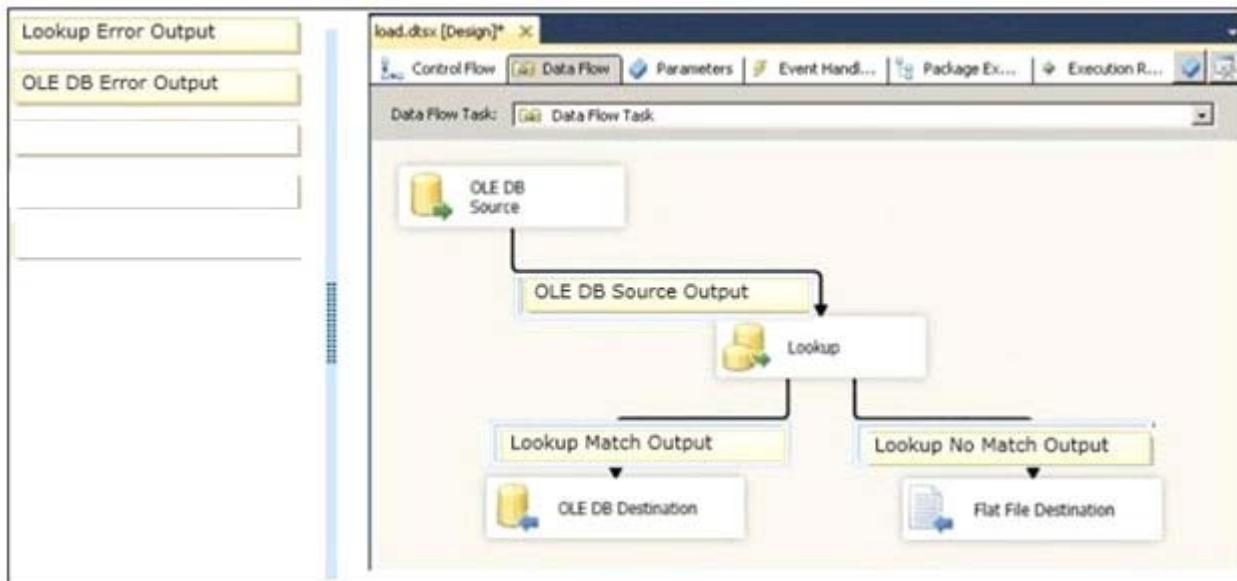
A.

**Correct Answer:**

**Section: (none)**

**Explanation**

**Explanation/Reference:**



### QUESTION 26

You are editing a SQL Server Integration Services (SSIS) package that contains three Execute SQL tasks and no other tasks. The three Execute SQL tasks modify products in staging tables in preparation for a data warehouse load. The package and all three Execute SQL product tasks have their TransactionOption property set to Supported. You need to ensure that if any of the three Execute SQL product tasks fail, all three tasks will roll back their changes. What should you do?

- Change the TransactionOption property of the package to Required.
- Change the TransactionOption property of all three Execute SQL product tasks to Required.
- Move the three Execute SQL product tasks into a Foreach Loop container.
- Move the three Execute SQL product tasks into a Sequence container.

**Correct Answer:** A

**Section:** (none)

**Explanation**

**Explanation/Reference:**

### QUESTION 27

You are developing a SQL Server Integration Services (SSIS) package that imports data into a data warehouse hosted on SQL Azure. The package uses a Foreach container to process text files found in a folder. The package must be deployed to a single server by using the Project Deployment model. Multiple SQL Server Agent jobs call the package. Each job is executed on a different schedule. Each job passes a different folder path to the package. You need to configure the package to accept the folder path from each job. Which package configuration should you use?

- .dtsConfig file
- Registry Entry
- Environment Variable
- Parent Package Variable
- XML Configuration File

**Correct Answer:** C

**Section:** (none)

## Explanation

### Explanation/Reference:

#### QUESTION 28

You are developing a SQL Server Integration Services (SSIS) package to load data into a data warehouse.



You need to establish the correct order for loading each table to maximize parallel processing. Which order should you use?

- A. 1. DimCurrency, DimScenario, DimAccount in parallel  
2. DimOrganization  
3. FactFinance
- B. 1. DimCurrency, DimOrganization in parallel  
2. DimScenario, DimAccount in parallel  
3. FactFinance
- C. 1. DimCurrency, FactFinance in parallel  
2. DimOrganization, DimScenario, DimAccount in parallel Get Complete Collection of 70-458 Exam's Question and Answers.
- D. 1. FactFinance  
2. DimOrganization, DimScenario, DimAccount in parallel  
3. DimCurrency
- E. 1. DimCurrency  
2. DimOrganization  
3. DimScenario, DimAccount in parallel  
4. FactFinance

**Correct Answer:** A

**Section:** (none)

**Explanation**

### Explanation/Reference:

#### QUESTION 29

You develop a SQL Server Integration Services (SSIS) package that imports SQL Azure data into a data warehouse every night. The SQL Azure data contains many misspellings and variations of abbreviations. To import the data, a developer used the Fuzzy Lookup transformation to choose the closest-matching string

from a reference table of allowed values. The number of rows in the reference table is very large. If no acceptable match is found, the Fuzzy Lookup transformation passes a null value. The current setting for the Fuzzy Lookup similarity threshold is 0.50. Many values are incorrectly matched. You need to ensure that more accurate matches are made by the Fuzzy Lookup transformation without degrading performance. What should you do?

- A. Change the Exhaustive property to True,
- B. Change the similarity threshold to 0.55.
- C. Change the similarity threshold to 0.40.
- D. Increase the maximum number of matches per lookup.

**Correct Answer:** B

**Section:** (none)

**Explanation**

**Explanation/Reference:**

### **QUESTION 30**

To facilitate the troubleshooting of SQL Server Integration Services (SSIS) packages, a logging methodology is put in place. The methodology has the following requirements:

- The deployment process must be simplified.
- All the logs must be centralized in SQL Server.
- Log data must be available via reports or T-SQL.
- Log archival must be automated.

You need to configure a logging methodology that meets the requirements while minimizing the amount of deployment and development effort. What should you do?

- A. Open a command prompt and run the gacutil command.
- B. Open a command prompt and execute the package by using the SQL Log provider and running the dtexecui.exe utility.
- C. Add an OnError event handler to the SSIS project.
- D. Use an msi file to deploy the package on the server.
- E. Configure the output of a component in the package data flow to use a data tap.
- F. Run the dtutil command to deploy the package to the SSIS catalog and store the configuration in SQL Server.
- G. Open a command prompt and run the dtexec /rep /conn command.
- H. Open a command prompt and run the dtutil /copy command.
- I. Open a command prompt and run the dtexec /dumperror /conn command.
- J. Configure the SSIS solution to use the Project Deployment Model.
- K. Create a reusable custom logging component and use it in the SSIS project.

**Correct Answer:** J

**Section:** (none)

**Explanation**

**Explanation/Reference:**

### **QUESTION 31**

You are deploying a new SQL Server Integration Services (SSIS) project to the test environment. A package in the project uses a custom task component. You need to ensure that the custom object is deployed on the test environment correctly. What should you do?

- A. Run the package by using the dtexec /rep /conn command.
- B. Create a reusable custom logging component.
- C. Create an OnError event handler.
- D. Use the gacutil command.
  - E Use the dtutil /copy command.
- E. Deploy the package to the Integration Services catalog by using dtutil and use SQL Server to store the configuration.
- F. Run the package by using the dtexec /dumperror /conn command.
- G. Use the Project Deployment Wizard.
- H. Deploy the package by using an msi file.
- I. Add a data tap on the output of a component in the package data flow.
- J. Run the package by using the dtexecui.exe utility and the SQL Log provider.

**Correct Answer:** D

**Section:** (none)

**Explanation**

**Explanation/Reference:**

### **QUESTION 32**

You are deploying a project to the SQL Server Integration Services (SSIS) catalog. You need to ensure that the project is deployed to the production environment with the least amount of administrative effort. What should you do?

- A. Create an On Error event handler.
- B. Add a data tap on the output of a component in the package data flow.
- C. Use the gacutil command.
- D. Create a reusable custom logging component.
- E. Run the package by using the dtexecui.exe utility and the SQL Log provider.
- F. Run the package by using the dtexec /rep /conn command.
- G. Use the Project Deployment Wizard.
- H. Run the package by using the dtexec /dumperror /conn command.
- I. Use the dtutil /copy command.
- J. Deploy the package to the Integration Services catalog by using dtutil and use SQL Server to store the configuration.
- K Deploy the package by using an msi file.

**Correct Answer:** G

**Section:** (none)

**Explanation**

**Explanation/Reference:**

### **QUESTION 33**

You are migrating a database named Orders to a new server that runs Microsoft SQL Server 2012. You attempt to add a SQL login, [User1], to the database. However, you receive the following error message: "User already exists in current database."

You need to configure the [User1] login to be able to access the Orders database and retain the original permissions. You need to achieve this goal by using the minimum required permissions.

Which Transact-SQL statement should you use?

- A. DROP USER [User1];  
CREATE USER [User1]FOR LOGIN [User1]; ALTER ROLE [db\_owner] ADD MEMBER [User1];
- B. ALTER ROLE ;db\_owner] ADD MEMBER [User1];
- C. ALTER USER [User1]WITH LOGIN [User1];
- D. ALTER SERVER ROLE [sysadmin] ADD MEMBER [User1];

**Correct Answer:** C

**Section:** (none)

**Explanation**

**Explanation/Reference:**

#### **QUESTION 34**

You use a contained database named ContosoDb within a domain. You need to create a user who can log on to the ContosoDb database. You also need to ensure that you can port the database to different database servers within the domain without additional user account configurations. Which type of user should you create?

- A. SQL user with login
- B. SQL user without login
- C. Domain user
- D. User mapped to an asymmetric key

**Correct Answer:** A

**Section:** (none)

**Explanation**

**Explanation/Reference:**

#### **QUESTION 35**

You are implementing a SQL Server Integration Services (SSIS) package that loads data hosted in a SQL Azure database into a data warehouse. The source system contains redundant or inconsistent data. When the package finds invalid data, the row containing the invalid data must be omitted but it must also be written to a text file for further analysis. You need to establish the best technique to log these invalid rows while keeping the amount of development effort to a minimum. What should you do?

- A. Add an OnError event handler to the SSIS project.
- B. Use an msi file to deploy the package on the server.
- C. Open a command prompt and run the gacutil command.
- D. Open a command prompt and run the dtutil /copy command.
- E. Open a command prompt and run the dtexec /rep /conn command.
- F. Open a command prompt and run the dtexec /dumperror /conn command.
- G. Open a command prompt and execute the package by using the SQL Log provider and running the dtexecui.exe utility.
- H. Create a reusable custom logging component and use it in the SSIS project.
- I. Configure the SSIS solution to use the Project Deployment Model.
- J. Configure the output of a component in the package data flow to use a data tap.
- K. Run the dtutil command to deploy the package to the SSIS catalog and store the configuration in SQL Server.

**Correct Answer:** J

**Section: (none)****Explanation****Explanation/Reference:****QUESTION 36**

To ease the debugging of packages, you standardize the SQL Server Integration Services (SSIS) package logging methodology. The methodology has the following requirements:

- Centralized logging in SQL Server
- Simple deployment
- Availability of log information through reports or T-SQL
- Automatic purge of older log entries
- Configurable log details

You need to configure a logging methodology that meets the requirements while minimizing the amount of deployment and development effort. What should you do?

- A. Use the gacutil command.
- B. Use the dtutil /copy command.
- C. Use the Project Deployment Wizard.
- D. Create an OnError event handler.
- E. Create a reusable custom logging component.
- F. Run the package by using the dtexec /rep /conn command.
- G. Run the package by using the dtexec /dumperror /conn command.
- H. Run the package by using the dtexecui.exe utility and the SQL Log provider.
- I. Add a data tap on the output of a component in the package data flow.
- J. Deploy the package by using an msi file.
- K. Deploy the package to the Integration Services catalog by using dtutil and use SQL Server to store the configuration.

**Correct Answer: C****Section: (none)****Explanation****Explanation/Reference:****QUESTION 37**

You are troubleshooting an existing SQL Server Integration Services (SSIS) package. On several occasions, the package execution does not finish and no data seems to have been transferred. You need to ensure that package logging occurs. Your solution must minimize deployment and development efforts. What should you do?

- A. Add an OnError event handler to the SSIS project.
- B. Use an msi file to deploy the package on the server.
- C. Open a command prompt and run the gacutil command.
- D. Open a command prompt and run the dtutil /copy command.
- E. Open a command prompt and run the dtexec /rep /conn command.
- F. Open a command prompt and run the dtexec /dumperror /conn command.
- G. Open a command prompt and execute the package by using the SQL Log provider and running the dtexecui.exe utility.
- H. Create a reusable custom logging component and use it in the SSIS project.
- I. Configure the SSIS solution to use the Project Deployment Model.
- J. Configure the output of a component in the package data flow to use a data tap.

- K. Run the dtutil command to deploy the package to the SSIS catalog and store the configuration in SQL Server.

**Correct Answer:** I

**Section:** (none)

**Explanation**

**Explanation/Reference:**

### **QUESTION 38**

You are developing a SQL Server Integration Services (SSIS) project that copies a large amount of rows from a SQL Azure database. The project uses the Package Deployment Model. This project is deployed to SQL Server on a test server. You need to ensure that the project is deployed to the SSIS catalog on the production server. What should you do?

- A. Add an OnError event handler to the SSIS project.
- B. Use an msi file to deploy the package on the server.
- C. Open a command prompt and run the gacutil command.
- D. Open a command prompt and run the dtutil /copy command.
- E. Open a command prompt and run the dtexec /rep /conn command.
- F. Open a command prompt and run the dtexec /dumperror /conn command.
- G. Open a command prompt and execute the package by using the SQL Log provider and running the dtexecui.exe utility.
- H. Create a reusable custom logging component and use it in the SSIS project.
- I. Configure the SSIS solution to use the Project Deployment Model.
- J. Configure the output of a component in the package data flow to use a data tap.
- K. Run the dtutil command to deploy the package to the SSIS catalog and store the configuration in SQL Server.

**Correct Answer:** I

**Section:** (none)

**Explanation**

**Explanation/Reference:**

### **QUESTION 39**

You are developing a SQL Server Integration Services (SSIS) package that imports data into a data warehouse. You add an Execute SQL task to the control flow. The task must execute a simple INSERT statement. The task has the following requirements:

- The INSERT statement must use the value of a string package variable. The variable name is StringVar.
- The Execute SQL task must use an OLE DB Connection Manager. In the Parameter Mapping tab of the Execute SQL task, StringVar has been added as the only parameter. You must configure the SQLStatement property of the Execute SQL task. Which SQL statement should you use?

- A. INSERT INTO dbo.Table (variablevalue) VALUES (\$Project::StringVar)
- B. INSERT INTO dbo.Table (variablevalue) VALUES (@StringVar)
- C. INSERT INTO dbo.Table (variablevalue) VALUES (\$Package::StringVar)
- D. INSERT INTO dbo.Table (variablevalue) VALUES (?)

**Correct Answer:** D

**Section:** (none)

**Explanation**

**Explanation/Reference:****QUESTION 40**

You are editing a SQL Server Integration Services (SSIS) package that contains three Execute SQL tasks and no other tasks. The package and all three Execute SQL tasks have their TransactionOption property set to Supported. You need to ensure that if any of the Execute SQL tasks fail, all three tasks will roll back their changes. What should you do?

- A. Change the TransactionOption property of the package to Required.
- B. Change the TransactionOption property of all three Execute SQL tasks to Required.
- C. Move the three Execute SQL tasks into a Sequence container.
- D. Move the three Execute SQL tasks into a Foreach Loop container.

**Correct Answer:** A**Section:** (none)**Explanation****Explanation/Reference:****QUESTION 41**

You are designing a SQL Server Integration Services (SSIS) package that uploads a file to a table named Orders in a SQL Azure database. The company's auditing policies have the following requirements:

- An entry must be written to a dedicated SQL Server log table named OrderLog.
- The entry must be written as soon as the file upload task completes. You need to meet the company's policy requirements. Which event handler should you use?

- A. OnWarning
- B. OnVariableValueChanged
- C. OnExecStatusChanged
- D. OnPostExecute

**Correct Answer:** D**Section:** (none)**Explanation****Explanation/Reference:****QUESTION 42**

You are developing a SQL Server Integration Services (SSIS) package to load data into a SQL Server table on Server A. The package includes a data flow and is executed on ServerB. The destination table has its own identity column. The destination data load has the following requirements:

- The identity values from the source table must be used.
- Default constraints on the destination table must be ignored.

· Batch size must be 100,000 rows.

You need to add a destination and configure it to meet the requirements. Which destination should you use?

- A. ADO NET Destination with Bulk Insert
- B. SQL Server Destination
- C. OLE DB Destination with Fast Load
- D. ADO NET Destination without Bulk Insert
- E. OLE DB Destination without Fast Load

**Correct Answer:** C

**Section:** (none)

**Explanation**

**Explanation/Reference:**

**QUESTION 43**

You are adding a new capability to several dozen SQL Server Integration Services (SSIS) packages. The new capability is not available as an SSIS task. Each package must be extended with the same new capability. You need to add the new capability to all the packages without copying the code between packages. What should you do?

- A. Use the Expression task.
- B. Use the Script component.
- C. Use the Script task.
- D. Develop a custom task.
- E. Develop a custom component.

**Correct Answer:** D

**Section:** (none)

**Explanation**

**Explanation/Reference:**

**QUESTION 44**

You are designing a SQL Server Integration Services (SSIS) data flow to load sales transactions from a source system into a data warehouse hosted on SQL Azure. One of the columns in the data source is named ProductCode. Some of the data to be loaded will reference products that need special processing logic in the data flow. You need to enable separate processing streams for a subset of rows based on the source product code. Which data flow transformation should you use?

- A. Source Assistant
- B. Audit
- C. Destination Assistant
- D. Script Component

**Correct Answer:** B

**Section:** (none)

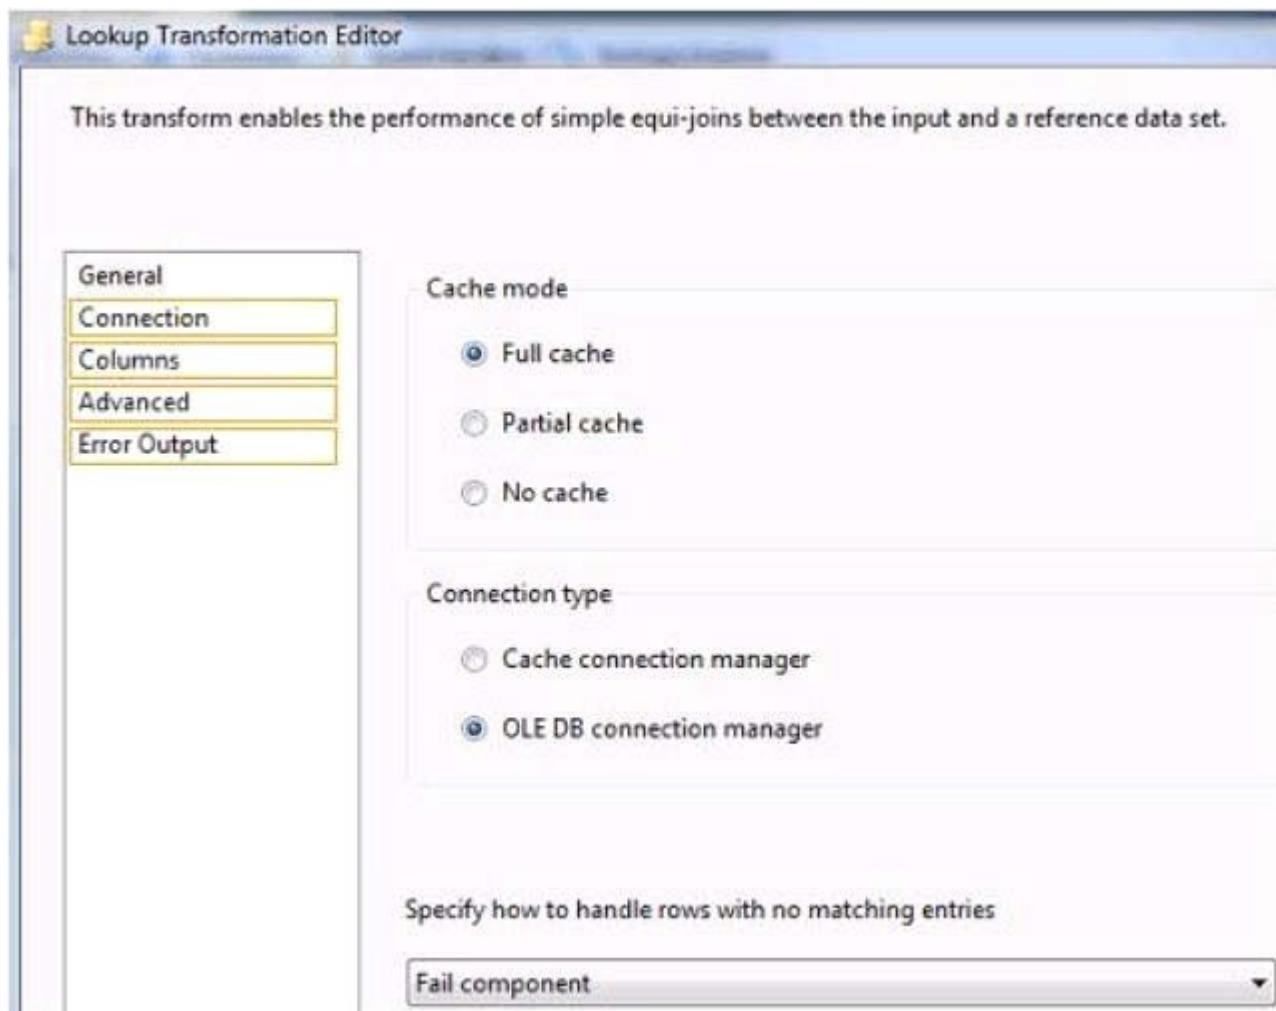
**Explanation**

**Explanation/Reference:**

**QUESTION 45**

**HOTSPOT**

You are developing a data flow to load sales data into a fact table. In the data flow, you configure a Lookup Transformation in full cache mode to look up the product data for the sale. The lookup source for the product data is contained in two tables. You need to set the data source for the lookup to be a query that combines the two tables. Which page of the Lookup Transformation Editor should you select to configure the query? To answer, select the appropriate page in the answer area.



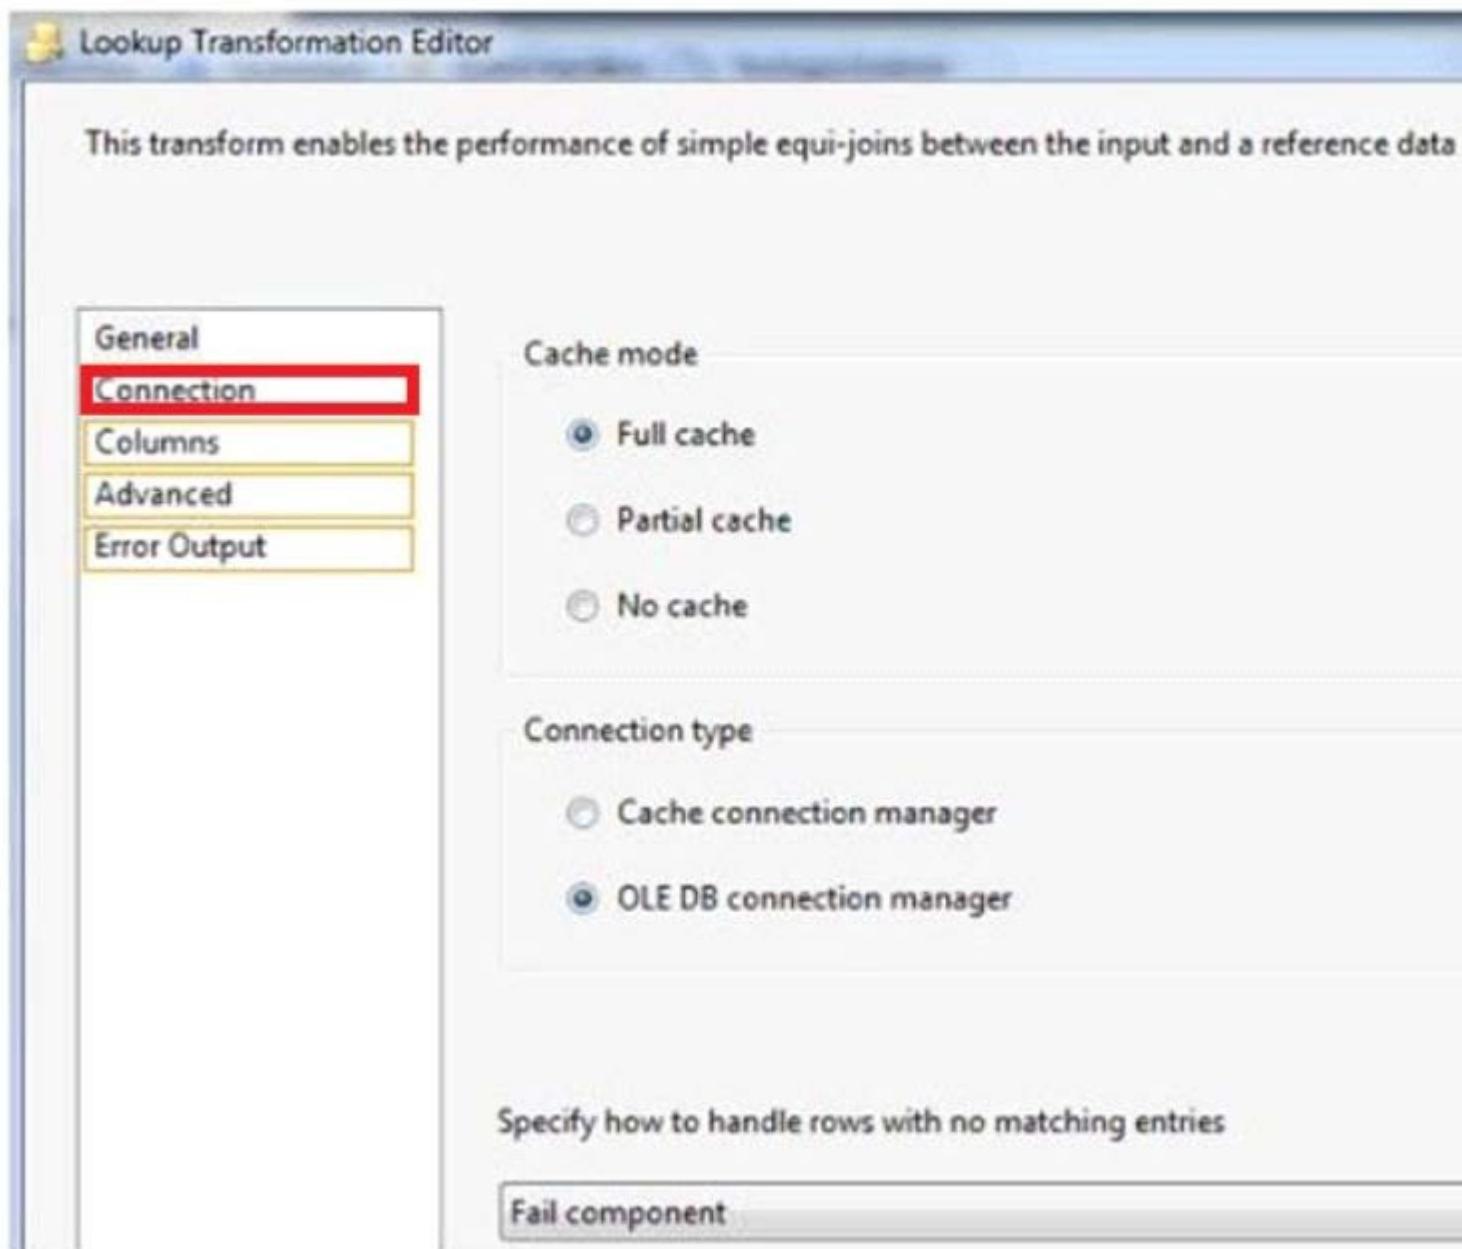
A.

**Correct Answer:**

**Section: (none)**

**Explanation**

**Explanation/Reference:**

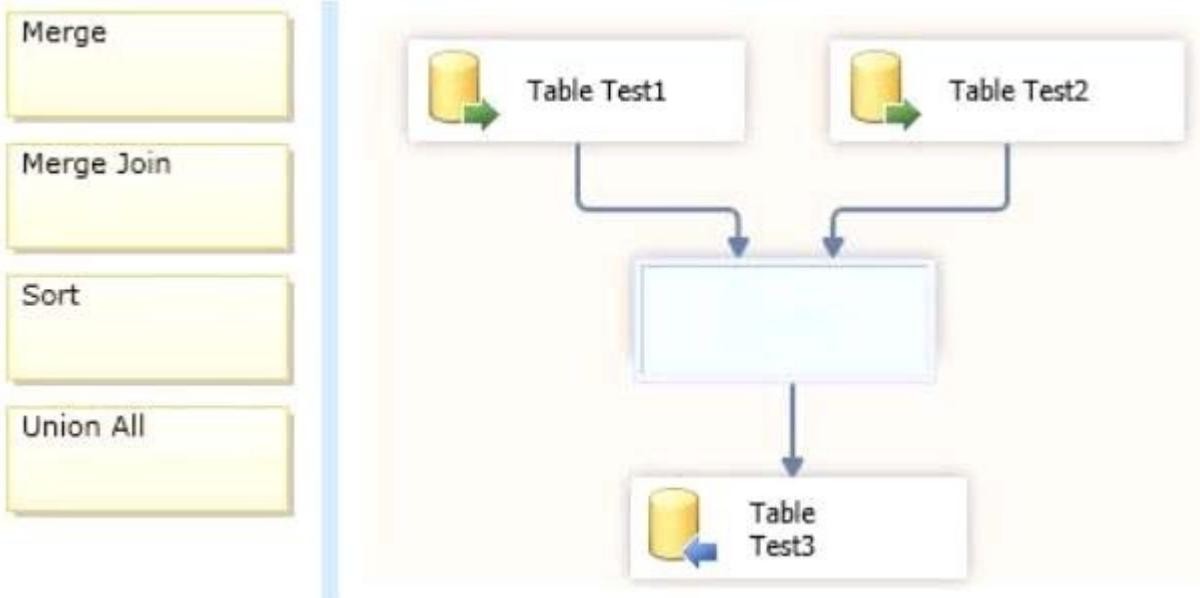


#### QUESTION 46

DRAG DROP

You are developing a SQL Server Integration Services (SSIS) package that imports unsorted data into a data warehouse hosted on SQL Azure. You have the following requirements:

- A destination table must contain all of the data in two source tables.
  - Duplicate records must be inserted into the destination table.
- You need to develop a data flow that imports the data while meeting the requirements. How should you develop the data flow? (To answer, drag the appropriate transformation from the list of transformations to the correct location in the answer area.)



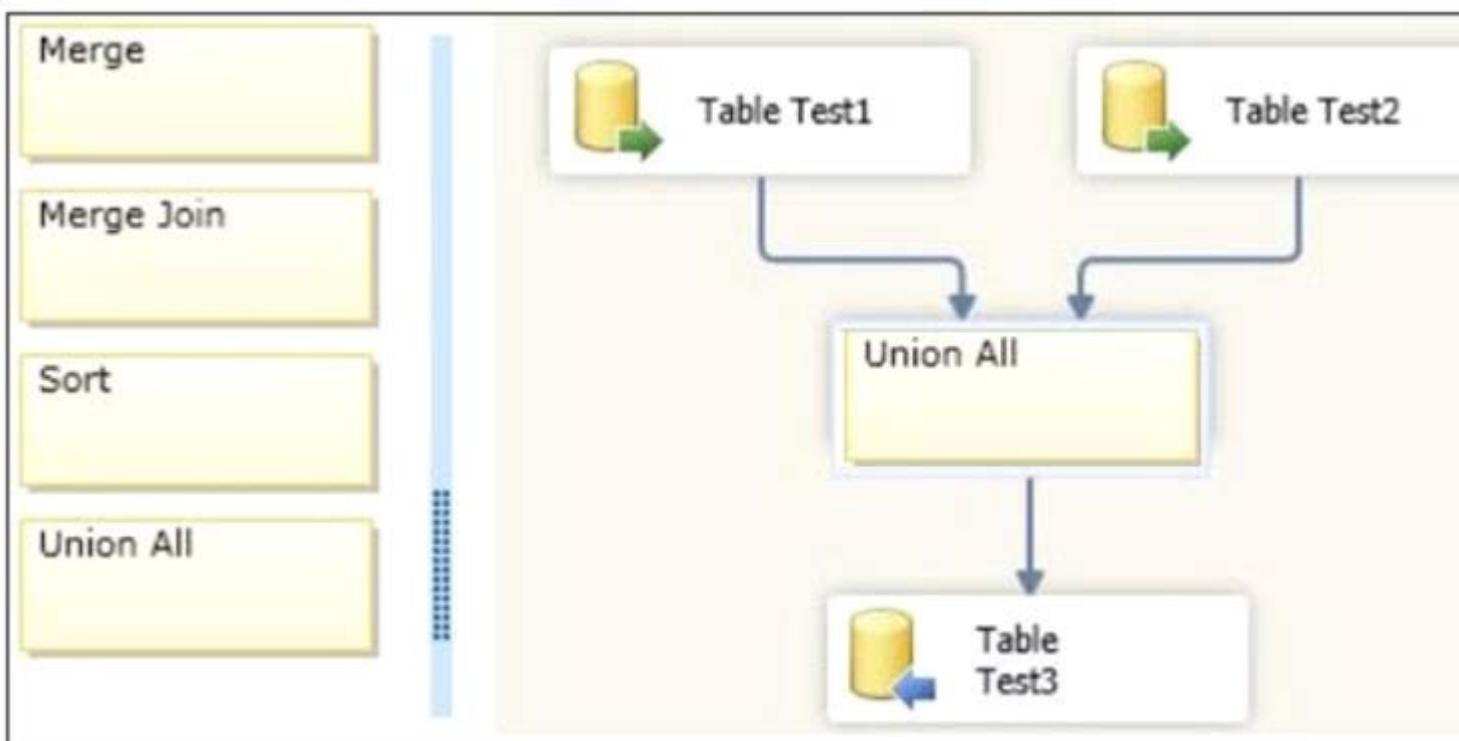
A.

**Correct Answer:**

**Section: (none)**

**Explanation**

**Explanation/Reference:**



#### QUESTION 47

You are creating a SQL Server Master Data Services (MDS) mode). This model is used to store a master list of products. An attribute must be added to the Product entity to define the sales manager responsible for each

product. You need to create an attribute in the Product entity that prevents users from entering invalid sales manager values. Which type of attribute should you create?

- A. Derived
- B. Domain-based
- C. User-defined
- D. Parent
- E. Explicit
- F. Recursive

**Correct Answer:** B

**Section:** (none)

**Explanation**

**Explanation/Reference:**

#### **QUESTION 48**

You develop a SQL Server Integration Services (SSIS) package that imports SQL Azure data into a data warehouse every night. The SQL Azure data contains many misspellings and variations of abbreviations. To import the data, a developer used the Fuzzy Lookup transformation to choose the closest-matching string from a reference table of allowed values. The number of rows in the reference table is very large. If no acceptable match is found, the Fuzzy Lookup transformation passes a null value. The current setting for the Fuzzy Lookup similarity threshold is 0.50. Many values are incorrectly matched. You need to ensure that more accurate matches are made by the Fuzzy Lookup transformation without degrading performance. What should you do?

- A. Decrease the maximum number of matches per lookup.
- B. Change the similarity threshold to 0.55.
- C. Change the Exhaustive property to True.
- D. Increase the maximum number of matches per lookup.

**Correct Answer:** B

**Section:** (none)

**Explanation**

**Explanation/Reference:**

#### **QUESTION 49**

You are designing a data warehouse hosted on SQL Azure. The data warehouse currently includes the dimUser and dimDistrict dimension tables and the factSales fact table. The dimUser table contains records for each user permitted to run reports against the warehouse, and the dimDistrict table contains information about sales districts. The system is accessed by users from certain districts, as well as by area supervisors and users from the corporate headquarters. You need to design a table structure to ensure that certain users can see sales data for only certain districts. Some users must be permitted to see sales data from multiple districts. What should you do?

- A. Create a userDistrict table that contains primary key columns from the dimUser and dimDistrict tables.
- B. Partition the factSales table on the district column.
- C. Add a district column to the dimUser table.
- D. For each district, create a view of the factSales table that includes a WHERE clause for the district.

**Correct Answer:** A

**Section:** (none)

## Explanation

Explanation/Reference:

### QUESTION 50

DRAG DROP

You are creating a sales data warehouse. When a product exists in the product dimension, you update the product name. When a product does not exist, you insert a new record. In the current implementation, the DimProduct table must be scanned twice, once for the insert and again for the update. As a result, inserts and updates to the DimProduct table take longer than expected. You need to create a solution that uses a single command to perform an update and an insert. How should you use a MERGE T-SQL statement to accomplish this goal? (To answer, drag the appropriate answer choice from the list of options to the correct location or locations in the answer area. You may need to drag the split bar between panes or scroll to view content.)

Source	
Target	
Product	
Inserted	
MATCHED THEN	
NOT MATCHED BY TARGET THEN	
NOT MATCHED BY PRODUCT AND (Target.ProductID is NULL) THEN	
MATCHED AND (Target.ProductID is NULL) THEN	
MERGE dbo.DimProduct AS	
USING (SELECT ProductID, ProductName, ProductColor, ProductCategory	
FROM dbo.StagingProduct AS	
ON (Target.ProductID = Source.ProductID)	
WHEN	
UPDATE SET Target.ProductName = Source.ProductName	
WHEN	
INSERT (ProductID, ProductName, ProductColor, ProductCategory)	
VALUES (Source.ProductID, Source.ProductName, Source.ProductColor, Source.ProductCategory)	
OUTPUT \$action, Inserted.*, Deleted.*;	

A.

Correct Answer:

Section: (none)

Explanation

### Explanation/Reference:

```
Product
Inserted
NOT MATCHED BY PRODUCT AND (Target.ProductID is NULL) THEN
MATCHED AND (Target.ProductID is NULL) THEN
MERGE dbo.DimProduct AS Target
USING (SELECT ProductID, ProductName, ProductColor, ProductCategory
      FROM dbo.StagingProduct) AS Source
ON (Target.ProductID = Source.ProductID)
WHEN MATCHED THEN
    UPDATE SET Target.ProductName = Source.ProductName
WHEN NOT MATCHED BY TARGET THEN
    INSERT (ProductID, ProductName, ProductColor, ProductCategory)
        VALUES (Source.ProductID, Source.ProductName, Source.ProductColor, Source.ProductCategory)
    OUTPUT $action, Inserted.*, Deleted.*;
```